

Stakeholder Engagement Plan

C10 Team Tana

1. Stakeholder Groups

For this clinical guidance SLM project, I will focus on **frontline and community health workers** and **patients and community members** in the low-resource settings our tool is designed to serve.

2. What Matters Most

Frontline and community health workers care most about clinical reliability and professional protection; they need guidance they can trust under time pressure, without exposing themselves to blame if a recommendation goes wrong, and without the tool undermining their own expertise or workflow. Patients and community members prioritize trust and dignity in their care, wanting assurance that any AI involvement is transparent and doesn't replace human attention or judgment. Both groups value a tool that is genuinely useful in real, resource-constrained conditions rather than one designed around assumptions from well-resourced healthcare systems that don't reflect their daily reality.

3. Potential Harms and Benefits

Potential harms: misplaced confidence in a fluent but factually mismatched answer, especially for emergency triage; erosion of health workers' own clinical judgment through over-reliance; patients receiving AI-mediated guidance without consent or awareness; exclusion of non-English speakers or oral-tradition communities whose health-seeking language isn't reflected in training data; reinforcing a two-tier system where only better-resourced facilities benefit.

Potential benefits: faster, more consistent guidance when specialist support isn't immediately available; a source-traceable "second opinion" that builds rather than replaces worker confidence; wider, more consistent access to evidence-based guidance across facilities; potential to narrow rather than widen the care gap our problem statement identifies, if deployed thoughtfully.

4. Selected Engagement Methods (Action Catalogue)

Based on the Action Catalogue, **Participatory Design workshops** with frontline health workers and **Deliberative Workshops** with patients and community members would be strong methods to ensure the tool is co-created, clinically sound, and reflective of community trust needs.

5. Why Each Method Fits

Frontline and Community Health Workers

Participatory Design is the right method for frontline and community health workers because they are not passive recipients of this tool, they are its direct operators, and every design decision about answer length, phrasing, and triage cues has immediate consequences for how safely and effectively they can do their jobs. Unlike a method built around one-off consultation or feedback collection, Participatory Design treats health workers as co-creators throughout the development process, which matters here because their expertise is precisely the kind that cannot be fully anticipated in advance by a team that has not worked frontline shifts themselves.

This method is suitable for this group for several concrete reasons. First, health workers possess situated, practical knowledge, they know what a real patient conversation sounds like, what phrasing causes confusion, and what triage cues actually change their behavior in the moment, that is best captured through hands-on engagement with real outputs rather than abstract discussion. Second, health workers have a direct stake in getting this right: an unsafe or unclear model output could expose them professionally, so their incentive to catch problems early is strong and well-aligned with our own goals. Third, this group is typically time-constrained, so a method that produces concrete, usable design changes from each session (rather than open-ended dialogue that may not converge on actionable outcomes) respects their limited availability while still yielding meaningful influence over the tool.

In practice, engagement would be conducted through a series of short, structured co-design sessions, ideally two to three iterations spaced across the model's development. In each session, we would present health workers with real sample outputs generated by the current model version, grounded in cases similar to what they encounter in their own work. Workers would be asked to flag any answer that feels unsafe, unclear, contradicts local protocol, or omits an important local consideration (for instance, medicine availability or referral pathways specific to their facility). Rather than simply collecting this feedback for our team to interpret independently, we would work through several examples together in real time, asking workers to suggest alternative phrasing or triage cues, and testing whether their suggested changes still fit within our answer-length and style constraints.

This approach also has a secondary benefit: it builds trust and buy-in among the health workers involved, since co-creation tends to produce a stronger sense of ownership than being asked to simply validate a finished product.

Patients and Community Members

Deliberative Workshops are the appropriate method for patients and community members because this group's central concerns are less about specific technical outputs and more about values, trust, consent, dignity, and comfort with AI involvement in something as personal as their own health care. These are exactly the kinds of concerns that surface through structured group discussion and deliberation, rather than through direct evaluation of model outputs, which would require a level of technical framing this group has no obligation or particular interest in engaging with.

This method suits patients and community members for several reasons specific to their position relative to the project. First, unlike health workers, this group does not interact with the tool directly, they experience its effects indirectly, through the guidance a health worker gives them. This means their most valuable input concerns the broader legitimacy and acceptability of the tool's use, not its specific wording, which is better explored through open, facilitated conversation than through granular design feedback. Second, deliberative formats do not require technical familiarity or literacy with AI concepts, which matters given this group may have widely varying levels of formal education and exposure to digital tools, a Deliberative Workshop's strength is in surfacing what people care about in their own language and frame of reference, rather than asking them to evaluate something in ours. Third, health-related trust is often deeply tied to community and cultural context rather than individual opinion, and a workshop format that brings community members together allows shared concerns, disagreements, and cultural considerations to surface collectively, which an individual interview approach would likely flatten or miss entirely.

In practice, this engagement would be conducted in partnership with a trusted local community health organization or existing community group structure, rather than being organized directly by our team, since trust in the facilitating body strongly shapes how openly participants engage on a topic as sensitive as health care. Each workshop would begin with a plain-language explanation of the tool's purpose and limitations explicitly avoiding overstating its capabilities followed by a small number of concrete example interactions read aloud or acted out, so participants can react to something tangible rather than an abstract description. Facilitators would then guide open discussion around specific prompts: how participants feel about a health worker using an AI-generated recommendation, what would make them trust or distrust such guidance, whether they would want to be told explicitly when AI was involved in their care, and whether anything about the example interactions felt culturally unfamiliar or inappropriate.

Because this is a deliberative rather than extractive format, the goal is not simply to gather opinions but to allow participants to reason through tensions together; for instance, weighing faster access to guidance against unfamiliarity with AI-mediated care producing more considered input than a quick survey would capture, and giving our team a clearer sense of the conditions under which this community would consider such a tool acceptable at all.

6. Influence on Dataset and Model Design

Health worker participation would directly refine our answer templates and triage phrasing, ensuring the model reflects real clinical workflow rather than assumptions made in isolation, and could surface locally missing question types to add to our dataset. Patient and community deliberation would shape consent and transparency practices around AI use in care, and could reveal cultural framing issues in how questions and answers are worded prompting revisions to tone or content before any real-world use. Together, both methods would help ensure the model is not just accurate by benchmark standards, but genuinely trusted and appropriate for its intended context.

References

[1] Kutoma Wakunuma and Tilimbe Jiya. 2019. Stakeholder Engagement and Responsible Research & Innovation in promoting Sustainable Development and Empowerment through ICT. *European Journal of Sustainable Development*, 8, 3, 275–275.

<https://doi.org/10.14207/ejsd.2019.v8n3p275>.

[2] Gendered Chatbots in Nigeria: Critical Perspectives. In *Responsible AI in Africa: Challenges and Opportunities*. Damian Okaibedi Eke, Kutoma Wakunuma, and Simisola Akintoye (Eds.). Springer International Publishing, 119–139.

https://doi.org/10.1007/978-3-031-08215-3_6

[3] The Engage2020 Action Catalogue. 2015. Retrieved from <https://actioncatalogue.eu/>.